



---

# CSPE65: MACHINE LEARNING TECHNIQUES AND PRACTICES

## Assignment - 2

---

106119100 - Rajneesh Pandey



# Problem Statements

1. Do the following for the given dataset:

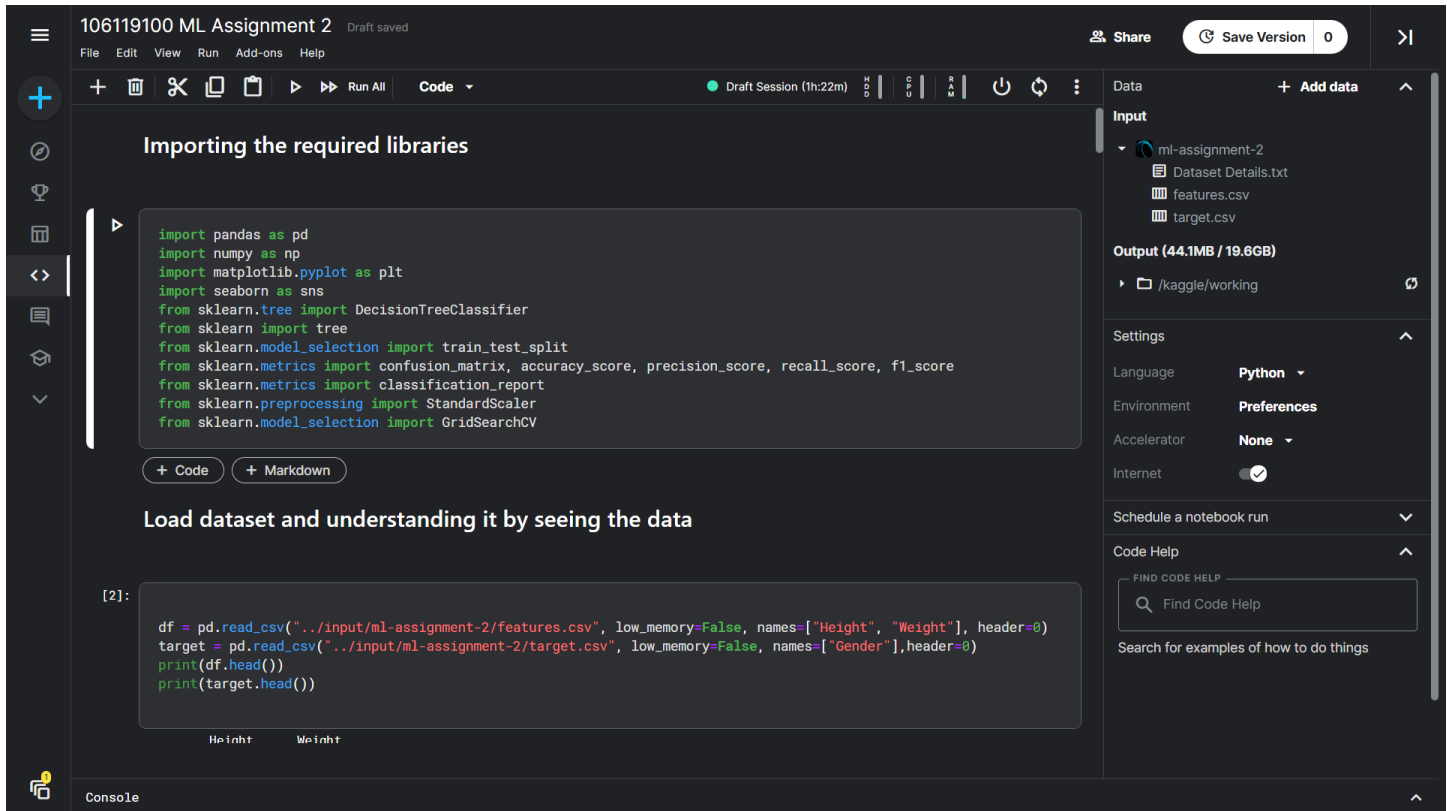
- a. Visualize the dataset using a plot (*Library: Seaborn and Matplotlib*) and create Decision Tree Algorithm to its complete depth (visualize the tree after construction).
  - Calculate the following evaluation metrics  $\rightarrow$  *Accuracy, Precision, Recall, F1 Score, Confusion Matrix* and discuss what you observe.
- b. Create Decision Tree Algorithm with hyperparameter tuning (visualize the tree after construction).
  - Calculate the following evaluation metrics  $\rightarrow$  *Accuracy, Precision, Recall, F1 Score, Confusion Matrix* and discuss what you observe.

2. Do the following for the given dataset:

- a. Apply k-Nearest Neighbour algorithm on the given dataset (*Find the best value for “k” using the method explained in class*).
  - Calculate the following evaluation metrics  $\rightarrow$  *Accuracy, Precision, Recall, F1 Score, Confusion Matrix* and discuss what you observe.
- b. Apply Min-Max Normalization on the given dataset and visualize it using a plot; Then, repeat Section “2. a” fully on the Normalized dataset.
- c. Plot ROC curves and calculate the corresponding AUC values for Section “2. a” and “2. b” and discuss what you observe.

# Code is written on Kaggle Notebook

Link of the notebook : <https://www.kaggle.com/rajneesh1708/106119100-ml-assignment-2/edit>

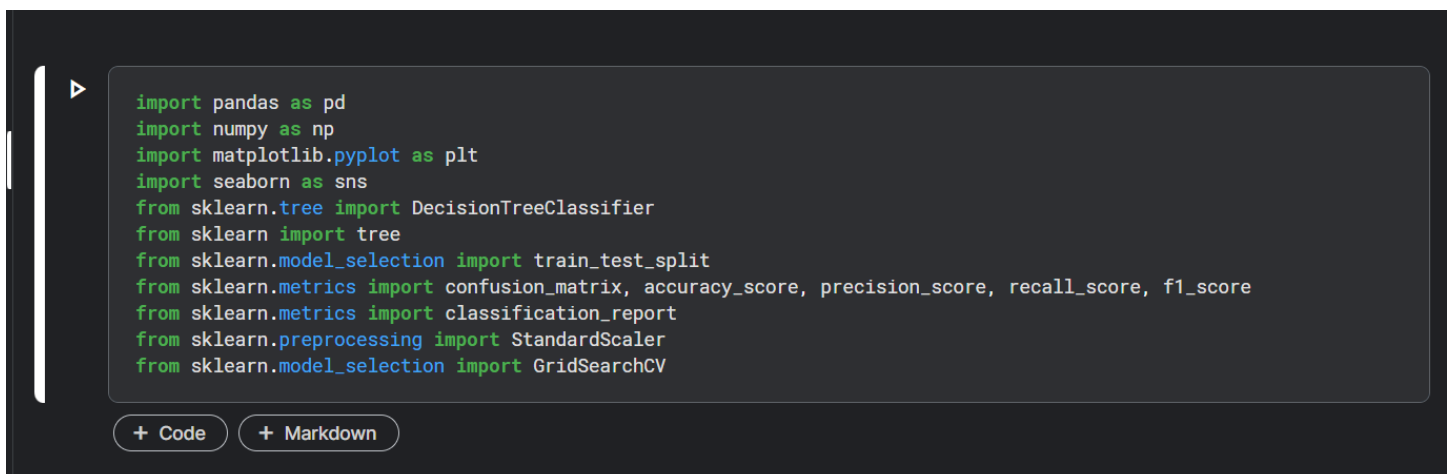


The screenshot shows a Kaggle Notebook titled "106119100 ML Assignment 2". The interface includes a top bar with "File", "Edit", "View", "Run", "Add-ons", and "Help" menus. A toolbar on the left contains icons for adding code, markdown, and running cells. The main area displays two code cells. The first cell, titled "Importing the required libraries", contains Python code to import various libraries. The second cell, titled "Load dataset and understanding it by seeing the data", contains code to load and inspect CSV files. The right sidebar shows the "Data" section with a list of files: "ml-assignment-2", "Dataset Details.txt", "features.csv", and "target.csv". Below this is the "Output" section showing the current directory as "/kaggle/working". The "Settings" section on the right allows configuration for Language (Python), Environment (Preferences), Accelerator (None), and Internet access (checked). A "Code Help" section is also visible at the bottom of the sidebar.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import classification_report
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

```
df = pd.read_csv("../input/ml-assignment-2/features.csv", low_memory=False, names=["Height", "Weight"], header=0)
target = pd.read_csv("../input/ml-assignment-2/target.csv", low_memory=False, names=["Gender"], header=0)
print(df.head())
print(target.head())
```

## Importing the required libraries



This is a close-up view of the first code cell from the notebook. It contains the same Python code for importing libraries as shown in the main screenshot. The code is as follows:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
from sklearn.metrics import classification_report
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

## Load dataset and understanding it by seeing the data

### Load dataset and understanding it by seeing the data

```
df = pd.read_csv("../input/ml-assignment-2/features.csv", low_memory=False, names=["Height", "Weight"], header=0)
target = pd.read_csv("../input/ml-assignment-2/target.csv", low_memory=False, names=["Gender"], header=0)
print(df.head())
print(target.head())
```

```
      Height      Weight
0  1.170866e+06  114.834354
1  1.326528e+06  52.048800
2  1.436353e+06  56.271637
3  1.667446e+06  112.761949
4  1.606816e+06  74.433175
Gender
0  Female
1  Female
2  Female
3   Male
4   Male
```

+ Code

+ Markdown

## Checking for rows duplicate

### Checking for rows duplicacy

```
[3]: any(df.duplicated())
```

```
[3]: False
```

Mapping the target variables, for Male as 1 and Female And checking the replacement done.

```
mapping_dict = {"Gender": {"Male": 1, "Female": 0}}
target.replace(mapping_dict, inplace=True)
print(target.dtypes)
target.head()
```

```
Gender    int64
dtype: object
```

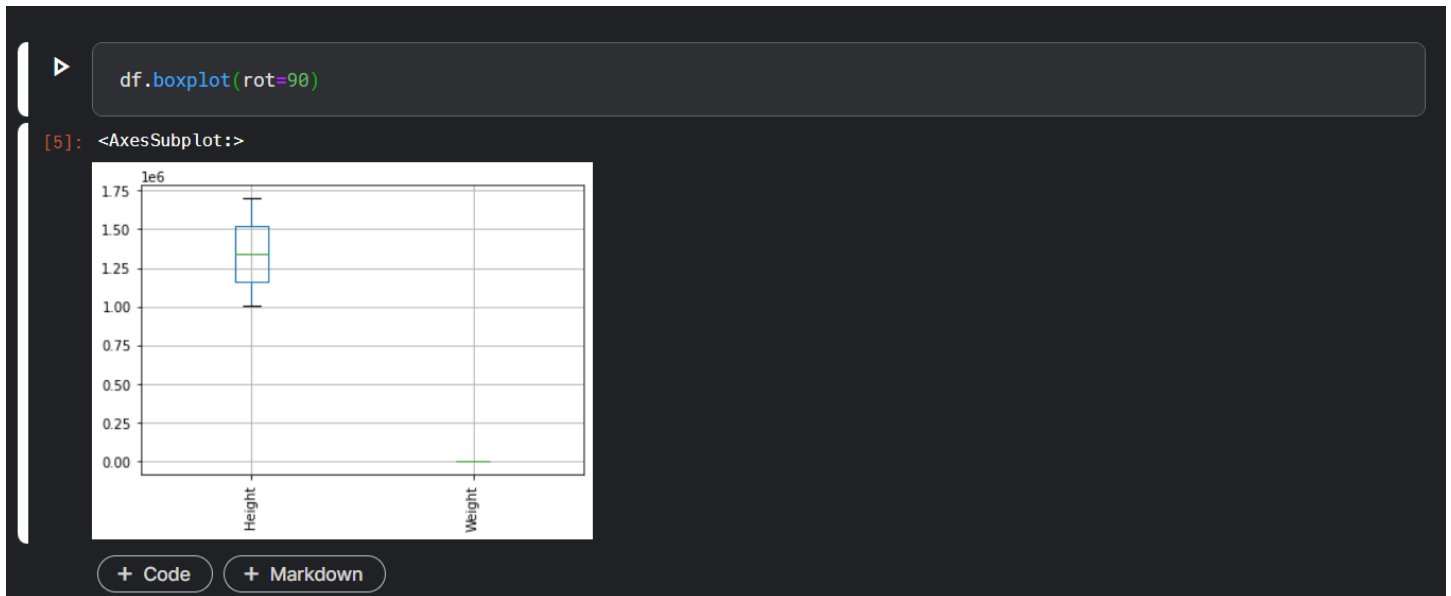
```
[4]: Gender
```

```
0    0
1    0
2    0
3    1
4    1
```

+ Code

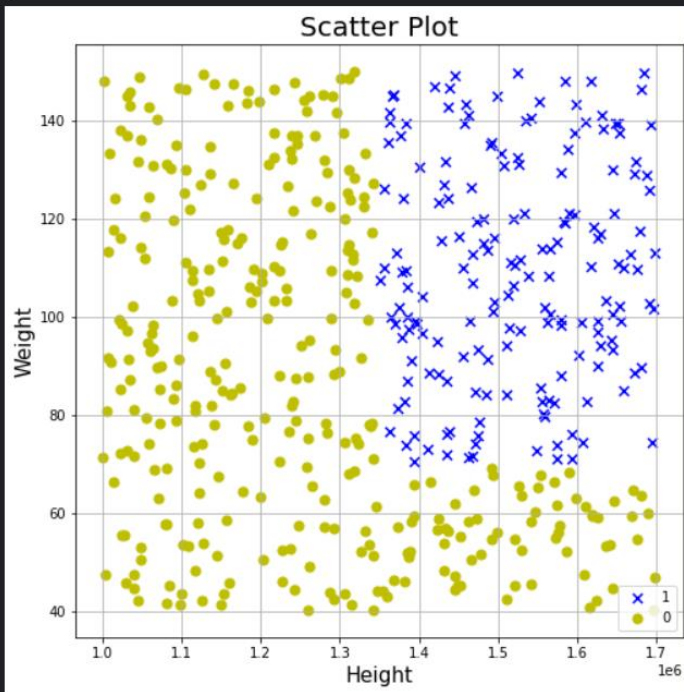
+ Markdown

## Checking the outliers present in the dataset by using Box and Whiskers Plot



## Scatter Plot

```
[6]: fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('Height', fontsize = 15)
ax.set_ylabel('Weight', fontsize = 15)
ax.set_title('Scatter Plot', fontsize = 20)
y_values = [1, 0]
colors = ['b', 'y']
markers = ['x', 'o']
for t, color, marker in zip(y_values, colors, markers):
    indicesToKeep = target['Gender'] == t
    ax.scatter(df.loc[indicesToKeep, 'Height'], df.loc[indicesToKeep, 'Weight'], c = color, s=50, marker=marker)
ax.legend(y_values)
ax.grid()
```

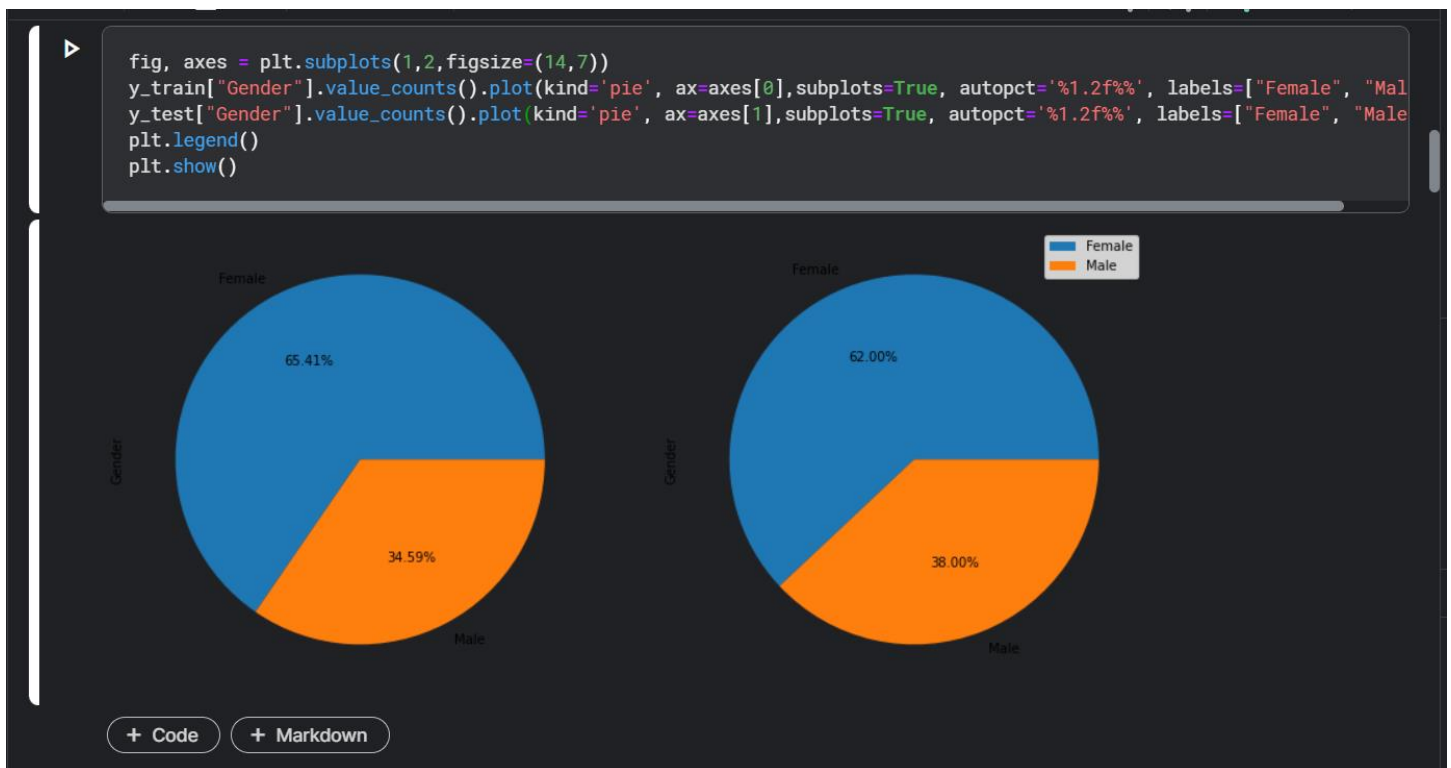


Split our dataset into training and testing dataset in the ratio of 80-20 respectively.

```
x_train, x_test, y_train, y_test = train_test_split(df, target, test_size=0.2, random_state=1)
print(len(x_train), len(x_test), len(y_train), len(y_test))
print(x_test.head(), y_test.head())
```

```
399 100 399 100
      Height      Weight
110  1.103653e+06  125.121096
147  1.521583e+06   45.075623
307  1.030528e+06   45.573982
326  1.427560e+06  115.348483
189  1.237990e+06   52.562252      Gender
110      0
147      0
307      0
326      1
189      0
```

Check for class imbalance i.e. we have data belonging to both the classes in training and testing dataset



As I can see from the plots, both the classes are split in appropriate proportions. So our model will not become biased after training.

## Question 1A and 1B

### Intialising the DecisionTreeClassifier

#### Question 1 A

+ Code

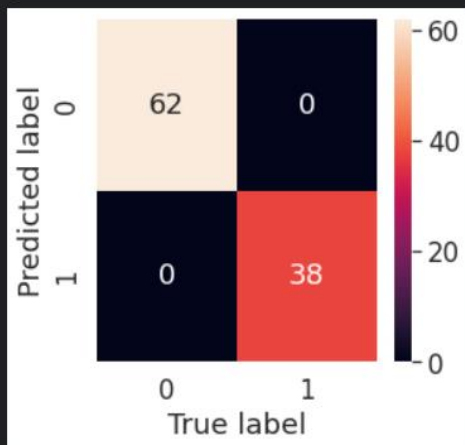
+ Markdown

```
[9]: model = DecisionTreeClassifier()
      fittedModel = model.fit(x_train, y_train)
      fittedModel
```

```
[9]: DecisionTreeClassifier()
```



```
predict = fittedModel.predict(x_test)
sns.set(font_scale=1.5)
fig, ax = plt.subplots(figsize=(4, 4))
ax = sns.heatmap(confusion_matrix(y_test, predict), annot=True, cbar=True)
plt.xlabel("True label")
plt.ylabel("Predicted label")
plt.show()
```



+ Code

+ Markdown



## Printing the required evaluation metrics (1A)

```
[11]: print("Confusion Matrix : \n",confusion_matrix(y_test, predict))
      print("Accuracy Score : ",accuracy_score(y_test, predict))
      print("Precision Score : ",precision_score(y_test, predict))
      print("Recall Score : ",recall_score(y_test, predict))
      print("F1 Score : ",f1_score(y_test, predict))
      target_names = ['Female', 'Male']
      print(classification_report(y_test, predict, target_names=target_names))
```

```
Confusion Matrix :
[[62  0]
 [ 0 38]]
Accuracy Score : 1.0
Precision Score : 1.0
Recall Score : 1.0
F1 Score : 1.0
```

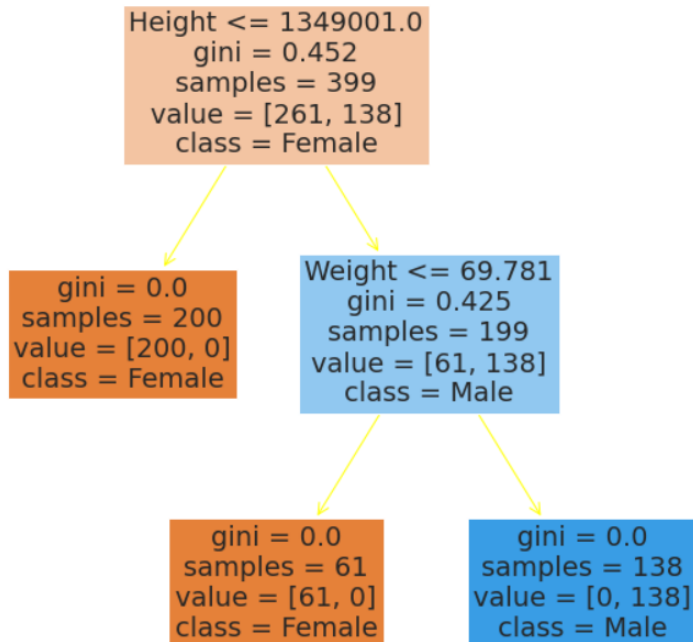
	precision	recall	f1-score	support
Female	1.00	1.00	1.00	62
Male	1.00	1.00	1.00	38
accuracy			1.00	100
macro avg	1.00	1.00	1.00	100
weighted avg	1.00	1.00	1.00	100

As we can see from the evaluation metrics, the model has an accuracy of 100% even without hyperparameter tuning. Since this is a small and easy dataset, we were able to get a 100% accuracy. But this need not be true for complex datasets.

```
fig = plt.figure(figsize=(10,10))
out = tree.plot_tree(fittedModel, feature_names=["Height", "Weight"], class_names=["Female", "Male"], filled=True)
for outs_ in out:
    arrow = outs_.arrow_patch
    if arrow is not None:
        arrow.set_edgecolor('yellow')
out
```

```
[12...] [Text(0.4, 0.8333333333333334, 'Height <= 1349001.0\ngini = 0.452\nsamples = 399\nvalue = [261, 138]\n\nclass = Female'),
Text(0.2, 0.5, 'gini = 0.0\nsamples = 200\nvalue = [200, 0]\n\nclass = Female'),
Text(0.6, 0.5, 'Weight <= 69.781\ngini = 0.425\nsamples = 199\nvalue = [61, 138]\n\nclass = Male'),
Text(0.4, 0.16666666666666666, 'gini = 0.0\nsamples = 61\nvalue = [61, 0]\n\nclass = Female'),
Text(0.8, 0.16666666666666666, 'gini = 0.0\nsamples = 138\nvalue = [0, 138]\n\nclass = Male')]
```

```
Text(0.6, 0.5, 'Weight <= 69.781\ngini = 0.425\nsamples = 199\nvalue = [61, 138]\nnclass = Male'),
Text(0.4, 0.16666666666666666, 'gini = 0.0\nsamples = 61\nvalue = [61, 0]\nnclass = Female'),
Text(0.8, 0.16666666666666666, 'gini = 0.0\nsamples = 138\nvalue = [0, 138]\nnclass = Male')]
```



This is the obtained Decision tree

gridsearchCV is used to tune the hyperparameters.

'min\_impurity\_decrease' is an efficient parameter to prevent overfitting of the model.

## Question 1 B

### Question 1 B

+ Code

+ Markdown

```
[13]: # RandomForestClassifier hyperparameters
grid_params = {"max_depth": [None, 1, 2, 4, 8], "min_samples_split": list(np.arange(2, 70, 10)),
               "min_samples_leaf": np.arange(1, 20, 2),
               "min_impurity_decrease": [0, 0.1, 0.2],
               "max_features": [1, 2]}
grid_model = GridSearchCV(DecisionTreeClassifier(), param_grid=grid_params, verbose=True)
grid_model.fit(x_train, y_train)
```

Fitting 5 folds for each of 2100 candidates, totalling 10500 fits

```
[13... GridSearchCV(estimator=DecisionTreeClassifier(),
                  param_grid={'max_depth': [None, 1, 2, 4, 8],
                              'max_features': [1, 2],
                              'min_impurity_decrease': [0, 0.1, 0.2],
                              'min_samples_leaf': array([ 1,  3,  5,  7,  9, 11, 13, 15, 17, 19]),
                              'min_samples_split': [2, 12, 22, 32, 42, 52, 62]},
                  verbose=True)
```

+ Code

+ Markdown

```
[14]: grid_model.best_params_
```

```
[14... {'max_depth': None,
       'max_features': 1,
       'min_impurity_decrease': 0,
       'min_samples_leaf': 5,
       'min_samples_split': 42}
```

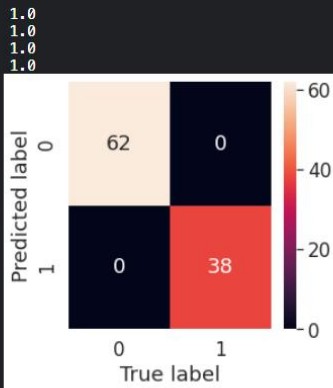
```
[15]: model = DecisionTreeClassifier(max_depth=grid_model.best_params_["max_depth"],
                                   max_features=grid_model.best_params_["max_features"],
                                   min_samples_leaf=grid_model.best_params_["min_samples_leaf"],
                                   min_impurity_decrease = grid_model.best_params_["min_impurity_decrease"],
                                   min_samples_split=grid_model.best_params_["min_samples_split"])
fittedModel = model.fit(x_train, y_train)
```

+ Code

+ Markdown

## Heat map

```
[16]: predict = fittedModel.predict(x_test)
sns.set(font_scale=1.5)
fig, ax = plt.subplots(figsize=(4, 4))
ax = sns.heatmap(confusion_matrix(y_test, predict), annot=True, cbar=True)
plt.xlabel("True label")
plt.ylabel("Predicted label")
print(accuracy_score(y_test, predict))
print(precision_score(y_test, predict))
print(recall_score(y_test, predict))
print(f1_score(y_test, predict))
plt.show()
```



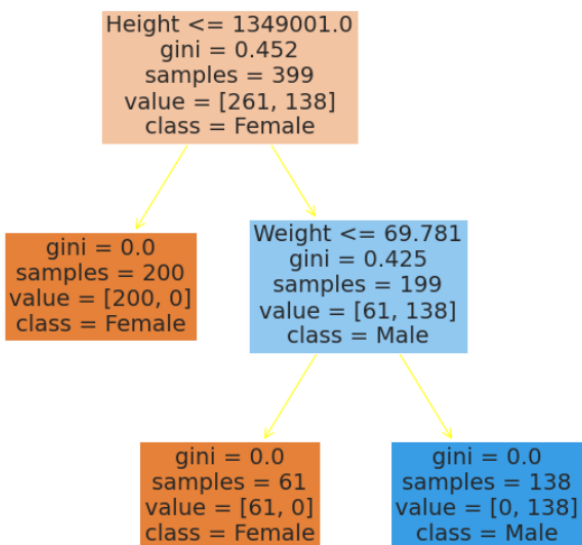
```
[17]: target_names = ['Female', 'Male']
print(classification_report(y_test, predict, target_names=target_names))
```

	precision	recall	f1-score	support
Female	1.00	1.00	1.00	62
Male	1.00	1.00	1.00	38
accuracy			1.00	100
macro avg	1.00	1.00	1.00	100
weighted avg	1.00	1.00	1.00	100

Here we got the same accuracy for both the trees.

But, for new complex real time data, this hyperparameter tuned decision tree may perform better

```
[18]: fig = plt.figure(figsize=(10,10))
out = tree.plot_tree(fittedModel, feature_names=["Height", "Weight"], class_names=["Female", "Male"], filled=True)
for o in out:
    arrow = o.arrow_patch
    if arrow is not None:
        arrow.set_edgecolor('yellow')
```



Reshape `y_train` and `y_test` to a one dimensional numpy array to suit KNN input format (avoid warnings)

```
[19]: y_train = y_train.values.ravel()
      y_train.shape
```

[19...] (399,)

```
[20]: y_test = y_test.values.ravel()
      y_test.shape
```

[20...] (100,)

## Question 2A

### Implementation of KNN

1. Import necessary class from sklearn
2. Initialise model
3. Fit model on training dataset

#### Question 2 A

```
[21]: from sklearn.neighbors import KNeighborsClassifier
      knn_model = KNeighborsClassifier(n_neighbors=5)
      knn_model.fit(x_train, y_train)
```

```
[21... KNeighborsClassifier()]
```

#### Make predictions on test split

[+ Code](#)[+ Markdown](#)

```
[22]: y_pred = knn_model.predict(x_test)
```

## Evaluating the model-confusion matrix and classification report

```
[23]: print(confusion_matrix(y_test, y_pred))
      print(classification_report(y_test, y_pred))
```

```
[[55  7]
 [ 3 35]]
```

	precision	recall	f1-score	support
0	0.95	0.89	0.92	62
1	0.83	0.92	0.88	38
accuracy			0.90	100
macro avg	0.89	0.90	0.90	100
weighted avg	0.90	0.90	0.90	100

[+ Code](#)[+ Markdown](#)

## Selecting best K value (number of nearest neighbors)

```
[24]: errors = []
      # Trying different K values
      for i in range(1, 60):
          knn = KNeighborsClassifier(n_neighbors=i)
          knn.fit(x_train, y_train)
          pred_i = knn.predict(x_test)
          errors.append(np.mean(pred_i != y_test))
```

+ Code

+ Markdown

```
[25]: #We print the values for each k value.
      for i in range (1,60):
          print('Error is',errors[i-1],', When k is',i)
```

Error is 0.17, When k is 1

+ ✂ 📄 📌 ▶ ▶▶ Run All

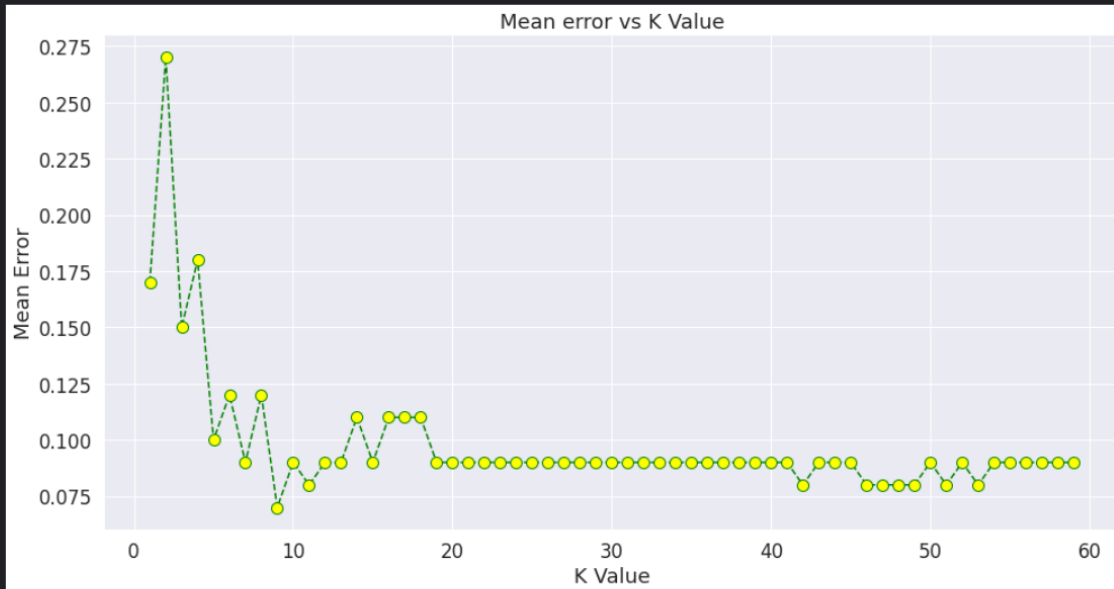
Markdown ▾

● Draft Session (1h:46m) H D C F A 🔌 ↺ ⋮

```
Error is 0.17, When k is 1
Error is 0.27, When k is 2
Error is 0.15, When k is 3
Error is 0.18, When k is 4
Error is 0.1, When k is 5
Error is 0.12, When k is 6
Error is 0.09, When k is 7
Error is 0.12, When k is 8
Error is 0.07, When k is 9
Error is 0.09, When k is 10
Error is 0.08, When k is 11
Error is 0.09, When k is 12
Error is 0.09, When k is 13
Error is 0.11, When k is 14
Error is 0.09, When k is 15
Error is 0.11, When k is 16
Error is 0.11, When k is 17
Error is 0.11, When k is 18
Error is 0.09, When k is 19
Error is 0.09, When k is 20
Error is 0.09, When k is 21
Error is 0.09, When k is 22
Error is 0.09, When k is 23
Error is 0.09, When k is 24
Error is 0.09, When k is 25
Error is 0.09, When k is 26
Error is 0.09, When k is 27
Error is 0.09, When k is 28
Error is 0.09, When k is 29
Error is 0.09, When k is 30
Error is 0.09, When k is 31
Error is 0.09, When k is 32
Error is 0.09, When k is 33
Error is 0.09, When k is 34
Error is 0.09, When k is 35
Error is 0.09, When k is 36
Error is 0.09, When k is 37
Error is 0.09, When k is 38
Error is 0.09, When k is 39
Error is 0.09, When k is 40
Error is 0.09, When k is 41
Error is 0.08, When k is 42
Error is 0.09, When k is 43
Error is 0.09, When k is 44
Error is 0.09, When k is 45
Error is 0.08, When k is 46
Error is 0.08, When k is 47
Error is 0.08, When k is 48
Error is 0.08, When k is 49
Error is 0.09, When k is 50
Error is 0.08, When k is 51
Error is 0.09, When k is 52
Error is 0.08, When k is 53
Error is 0.09, When k is 54
Error is 0.09, When k is 55
Error is 0.09, When k is 56
Error is 0.09, When k is 57
Error is 0.09, When k is 58
Error is 0.09, When k is 59
```

```
[26]: plt.figure(figsize=(16, 8))
plt.plot(range(1, 60), errors, color='green', linestyle='dashed', marker='o', markerfacecolor='yellow', markersize=10)
plt.title('Mean error vs K Value')
plt.xlabel('K Value')
plt.ylabel('Mean Error')
```

[26... Text(0, 0.5, 'Mean Error')



Creating a new model with n\_neighbors parameter as 25

```
[27]: knn_model = KNeighborsClassifier(n_neighbors=25)
knn_model.fit(x_train, y_train)
y_pred = knn_model.predict(x_test)
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.85	0.92	62
1	0.81	1.00	0.89	38
accuracy			0.91	100
macro avg	0.90	0.93	0.91	100
weighted avg	0.93	0.91	0.91	100

+ Code

+ Markdown

Confusion matrix:




## Confusion matrix:

9 labels that are predicted as 0 but true label is 1

```
[28]: sns.set(font_scale=1.5)
fig, ax = plt.subplots(figsize=(4, 4))
ax = sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cbar=True)
plt.xlabel("True label")
plt.ylabel("Predicted label")
```

```
[28... Text(4.5, 0.5, 'Predicted label')
```

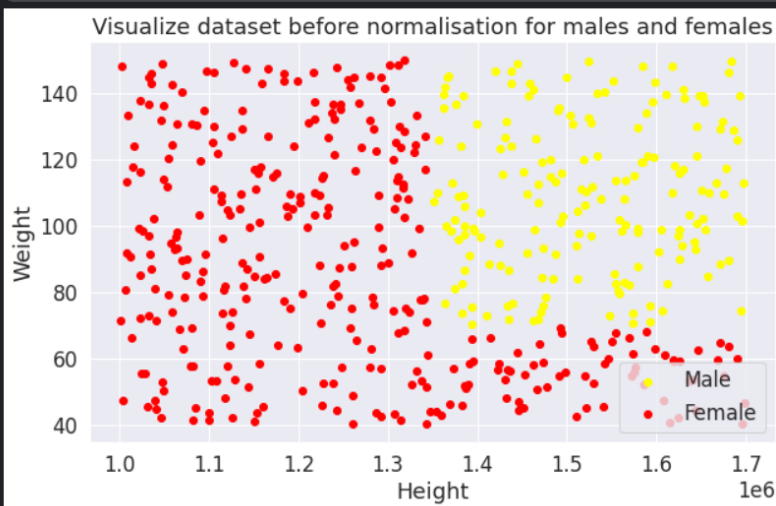


+ Code

+ Markdown

Visualising the dataset (data points under each target class) before normalization

```
[29]: df_out = pd.merge(df, target['Gender'], how = 'right', left_index = True,
right_index = True)
plt.figure(figsize=(10, 6))
plt.scatter(df_out.Height[df_out.Gender==1], df_out.Weight[df_out.Gender==1], c="yellow")
plt.scatter(df_out.Height[df_out.Gender==0], df_out.Weight[df_out.Gender==0], c="red")
plt.title("Visualize dataset before normalisation for males and females")
plt.xlabel("Height")
plt.ylabel("Weight")
plt.legend(["Male", "Female"]);
```



+ Code

+ Markdown

```
[30]: x_train.head()
```

```
[30...      Height  Weight
437  1.209246e+06  99.596811
273  1.431273e+06  53.645809
58   1.126067e+06  74.055455
379  1.365429e+06  99.802530
78   1.356239e+06  125.954853
```

[+ Code](#)[+ Markdown](#)

## Question 2 B

Min-max normalisation

Use same scalar we used for the train split to transform the test split

```
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaled_train = scaler.fit_transform(x_train)
scaled_test = scaler.transform(x_test)
```

[+ Code](#)[+ Markdown](#)

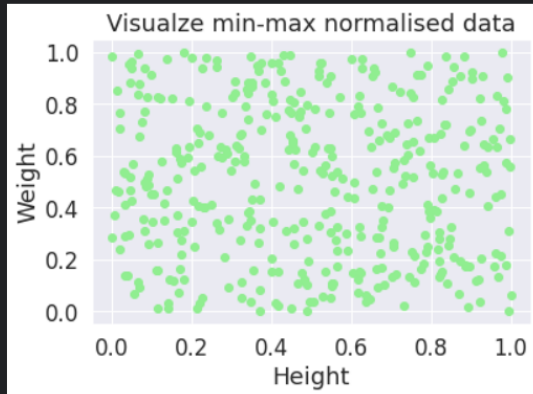
```
[32]: scaled_train[:5]
```

```
[32... array([[0.29818431, 0.54274625],
          [0.61676671, 0.12328948],
          [0.17883253, 0.30959585],
          [0.52228851, 0.54462412],
          [0.50910181, 0.78335165]])
```

## Visualizing data points after normalization

```
fig = plt.figure()
ax = fig.add_subplot(1, 1, 1)
x = scaled_train[:,0]
y = scaled_train[:,1]
ax.scatter(x, y, color = 'lightgreen')
plt.title("Visualze min-max normalised data")
plt.xlabel("Height")
plt.ylabel("Weight")
```

[33... Text(0, 0.5, 'Weight')



+ Code


+ Markdown

finding the optimal K value by plotting the error vs K value. But this time I use the normalised dataset to train the model

```
[34]: errors = []
# Trying different K values
for i in range(1, 60):
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(scaled_train, y_train)
    pred_i = knn.predict(scaled_test)
    errors.append(np.mean(pred_i != y_test))
```

```
plt.figure(figsize=(16, 8))
plt.plot(range(1, 60), errors, color='green', linestyle='dashed', marker='o', markerfacecolor='yellow', markersize=10)
plt.title('Error vs K')
plt.xlabel('K Value')
plt.ylabel('Average Error')
```

[35...] Text(0, 0.5, 'Average Error')



+ Code

+ Markdown

```
[36]: #We print the values for each k value.  
for i in range(1,60):  
    print('Error is',errors[i-1],', When k is',i)
```

```
Error is 0.0 , When k is 1  
Error is 0.0 , When k is 2  
Error is 0.0 , When k is 3  
Error is 0.0 , When k is 4  
Error is 0.0 , When k is 5  
Error is 0.0 , When k is 6  
Error is 0.0 , When k is 7  
Error is 0.0 , When k is 8  
Error is 0.02 , When k is 9  
Error is 0.01 , When k is 10  
Error is 0.01 , When k is 11  
Error is 0.01 , When k is 12  
Error is 0.02 , When k is 13  
Error is 0.01 , When k is 14  
Error is 0.0 , When k is 15  
Error is 0.01 , When k is 16  
Error is 0.01 , When k is 17  
Error is 0.0 , When k is 18  
Error is 0.0 , When k is 19  
Error is 0.01 , When k is 20  
Error is 0.0 , When k is 21  
Error is 0.01 , When k is 22  
Error is 0.0 , When k is 23  
Error is 0.01 , When k is 24  
Error is 0.0 , When k is 25  
Error is 0.02 , When k is 26  
Error is 0.02 , When k is 27  
Error is 0.02 , When k is 28  
Error is 0.02 , When k is 29  
Error is 0.02 , When k is 30  
Error is 0.02 , When k is 31  
Error is 0.02 , When k is 32  
Error is 0.02 , When k is 33  
Error is 0.02 , When k is 34  
Error is 0.01 , When k is 35  
Error is 0.01 , When k is 36  
Error is 0.01 , When k is 37  
Error is 0.02 , When k is 38  
Error is 0.02 , When k is 39  
Error is 0.02 , When k is 40  
Error is 0.02 , When k is 41  
Error is 0.02 , When k is 42  
Error is 0.02 , When k is 43  
Error is 0.02 , When k is 44  
Error is 0.02 , When k is 45  
Error is 0.02 , When k is 46  
Error is 0.02 , When k is 47  
Error is 0.02 , When k is 48  
Error is 0.02 , When k is 49  
Error is 0.02 , When k is 50  
Error is 0.02 , When k is 51  
Error is 0.02 , When k is 52  
Error is 0.02 , When k is 53  
Error is 0.02 , When k is 54  
Error is 0.02 , When k is 55  
Error is 0.02 , When k is 56  
Error is 0.02 , When k is 57  
Error is 0.02 , When k is 58  
Error is 0.02 , When k is 59
```

+ Code

+ Markdown

K=5 is the optimal value as obtained from the above graph.  
Now I will create a new model with n\_neighbors tuned to 5

```
[37]: knn_model = KNeighborsClassifier(n_neighbors=5)
      knn_model.fit(scaled_train, y_train)
```

```
[37... KNeighborsClassifier()
```

## Prediction on the test dataset

```
[38]: y_pred = knn_model.predict(scaled_test)
```

[+ Code](#)[+ Markdown](#)


## Classification Report and confusion matrix of the predicted values

```
[39]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	62
1	1.00	1.00	1.00	38
accuracy			1.00	100
macro avg	1.00	1.00	1.00	100
weighted avg	1.00	1.00	1.00	100

```
[40]: sns.set(font_scale=1.5)
      fig, ax = plt.subplots(figsize=(4, 4))
      ax = sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cbar=True)
      plt.xlabel("True label")
      plt.ylabel("Predicted label")
```

```
[40... Text(4.5, 0.5, 'Predicted label')
```



Again Using gridsearch to find the best parameters (like for n\_neighbors)



```
params = {'n_neighbors': np.arange(5, 60, 5)}  
gs = GridSearchCV(KNeighborsClassifier(), param_grid=params, cv=5, verbose=True)  
gs.fit(scaled_train, y_train)
```

Fitting 5 folds for each of 11 candidates, totalling 55 fits

```
[41...] GridSearchCV(cv=5, estimator=KNeighborsClassifier(),  
             param_grid={'n_neighbors': array([ 5, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55])},  
             verbose=True)
```

+ Code

+ Markdown

Got best value for n\_neighbors as 5 which matches with the same value obtained from the K vs error plot above



```
gs.best_params_
```

```
[42...] {'n_neighbors': 5}
```

+ Code

+ Markdown

Accuracy of 100% on the test split data is achieved

[43]:

```
gs.score(scaled_test, y_test)
```

```
[43...] 1.0
```

## Question 2C

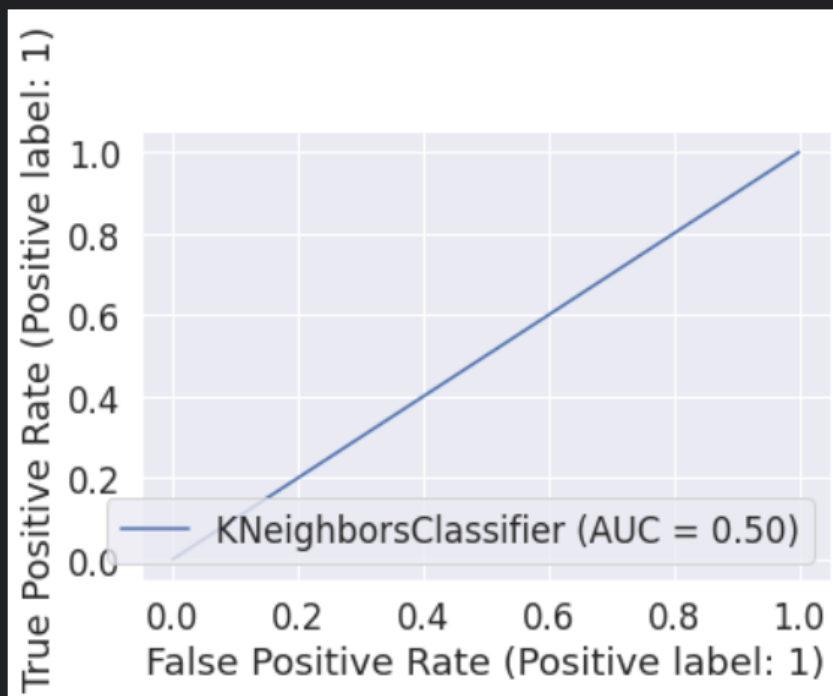
### ROC curves and AUC

#### Question 2 C

```
[44]: from sklearn.metrics import plot_roc_curve
```

```
[45]: model_diff_neighbors = KNeighborsClassifier(n_neighbors=5)
      model_diff_neighbors.fit(x_train, y_train)
      plot_roc_curve(model_diff_neighbors, scaled_test, y_test)
```

```
[45]... <sklearn.metrics._plot.roc_curve.RocCurveDisplay at 0x7fca64a9f550>
```



+ Code

+ Markdown

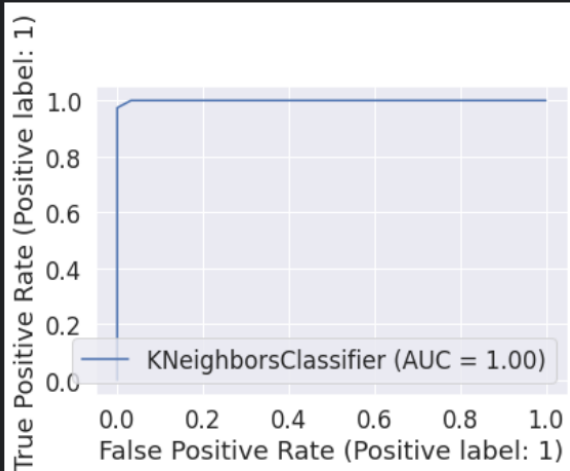
Got slightly bent roc curve (AUC not exactly one - rounded off in the plot) if I use different number of `n_neighbors` like 11 as the best hyperparameter is 5 for `n_neighbors`



```
[46]: model_diff_neighbors = KNeighborsClassifier(n_neighbors=9)
      model_diff_neighbors.fit(scaled_train, y_train)
      plot_roc_curve(model_diff_neighbors, scaled_test, y_test)
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_roc_curve
is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the
class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDispla
y.from_estimator`.
  warnings.warn(msg, category=FutureWarning)
```

```
[46... <sklearn.metrics._plot_roc_curve.RocCurveDisplay at 0x7fca64a66510>
```




Perfect ROC curve

Perfect plot with well tuned hyperparameter( $n\_neighbors=5$ ) and fitted on normalised data. AUC is one, hence this is the most optimum solution

```
[47]: plot_roc_curve(gs, scaled_test, y_test)
```

```
/opt/conda/lib/python3.7/site-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_roc_curve
is deprecated; Function :func:`plot_roc_curve` is deprecated in 1.0 and will be removed in 1.2. Use one of the
class methods: :meth:`sklearn.metrics.RocCurveDisplay.from_predictions` or :meth:`sklearn.metrics.RocCurveDispla
y.from_estimator`.
  warnings.warn(msg, category=FutureWarning)
```

```
[47... <sklearn.metrics._plot_roc_curve.RocCurveDisplay at 0x7fca649f8f90>
```



+ Code

+ Markdown

# Thankyou